

Neural Networks Training, SGD and Backpropagation

Machine Learning Course - CS-433

Nov 9, 2022

Nicolas Flammarion

EPFL

Recap

NNs: Key Facts

Supervised learning : we observe some data $S_{\text{train}} = \{x_n, y_n\}_{n=1}^N \in \mathcal{X} \times \mathcal{Y}$

➡ given a new x , we want to predict its label y

Linear prediction (with augmented features): $y = f_{\text{Lin}}(x) = \phi(x)^{\top} w$
Features are given

Prediction with a NN:

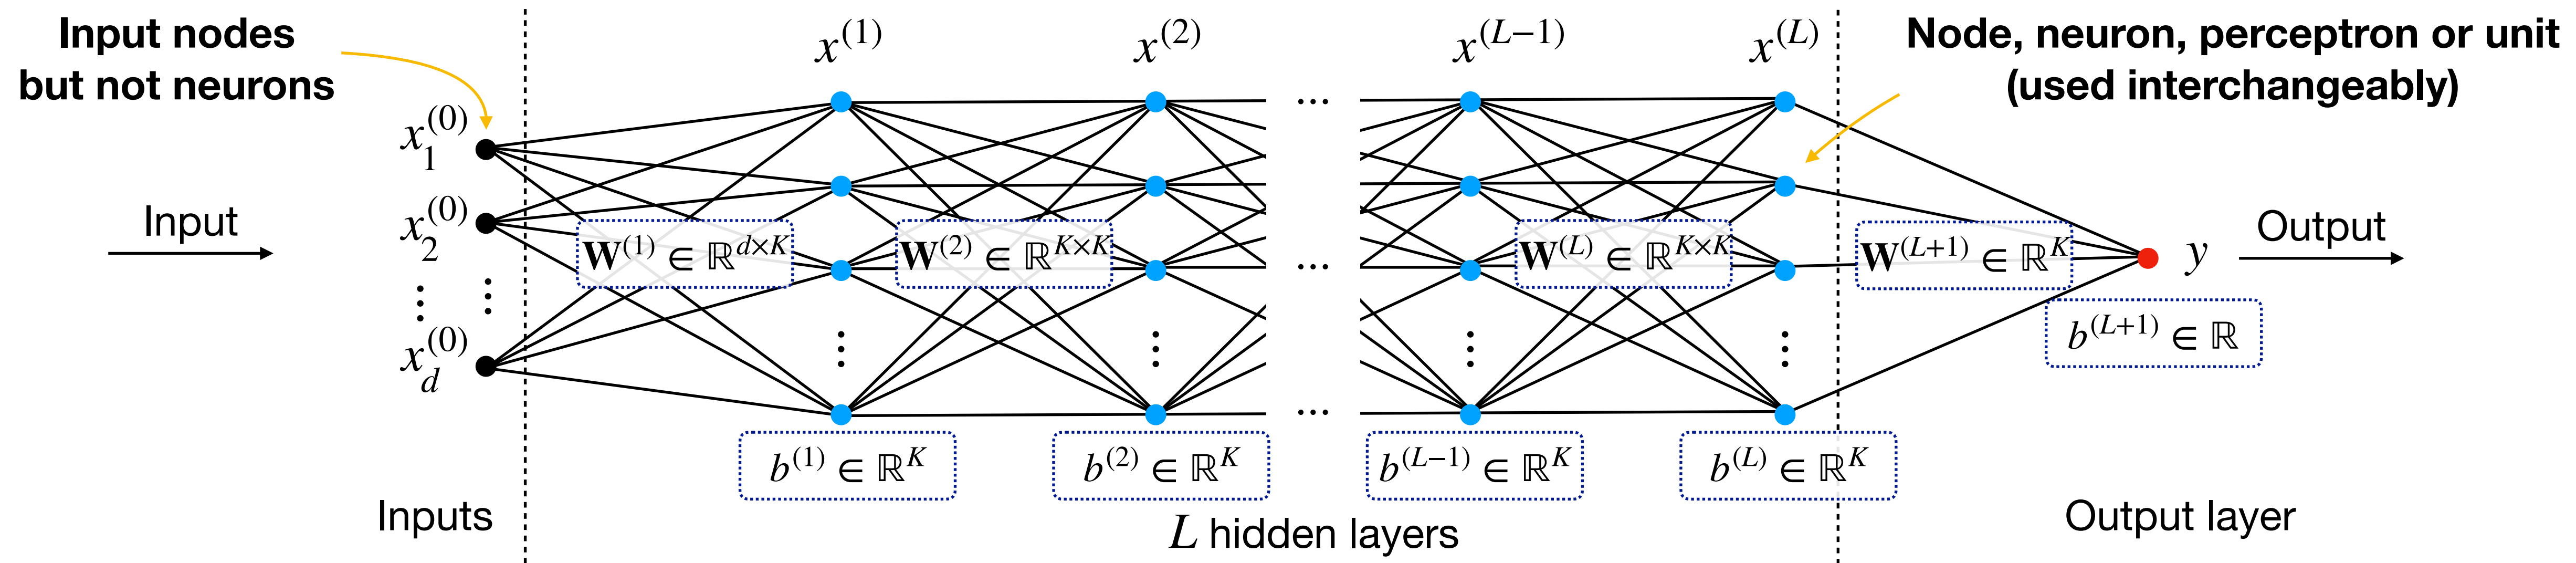
$y = f_{\text{NN}}(x) = f(x)^{\top} w$

Function implemented by the NN
parameters: weights and biases

First layers transform the input
into a good representation

Last layer is performing
a linear prediction

Fully Connected Neural Networks



Assume L hidden layers with K neurons each + output layer with single node


Outputs of hidden layer l given by vector: $x^{(l)} = f^{(l)}(x^{(l-1)}) := \phi((\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)})$

Elementwise

Learnable Parameters: Weight matrices $\mathbf{W}^{(l)}$ and bias vectors $b^{(l)}$ for $1 \leq l \leq L + 1$

- Each column of $\mathbf{W}^{(l)}$ corresponds to the weights of one perceptron

Single Neuron View



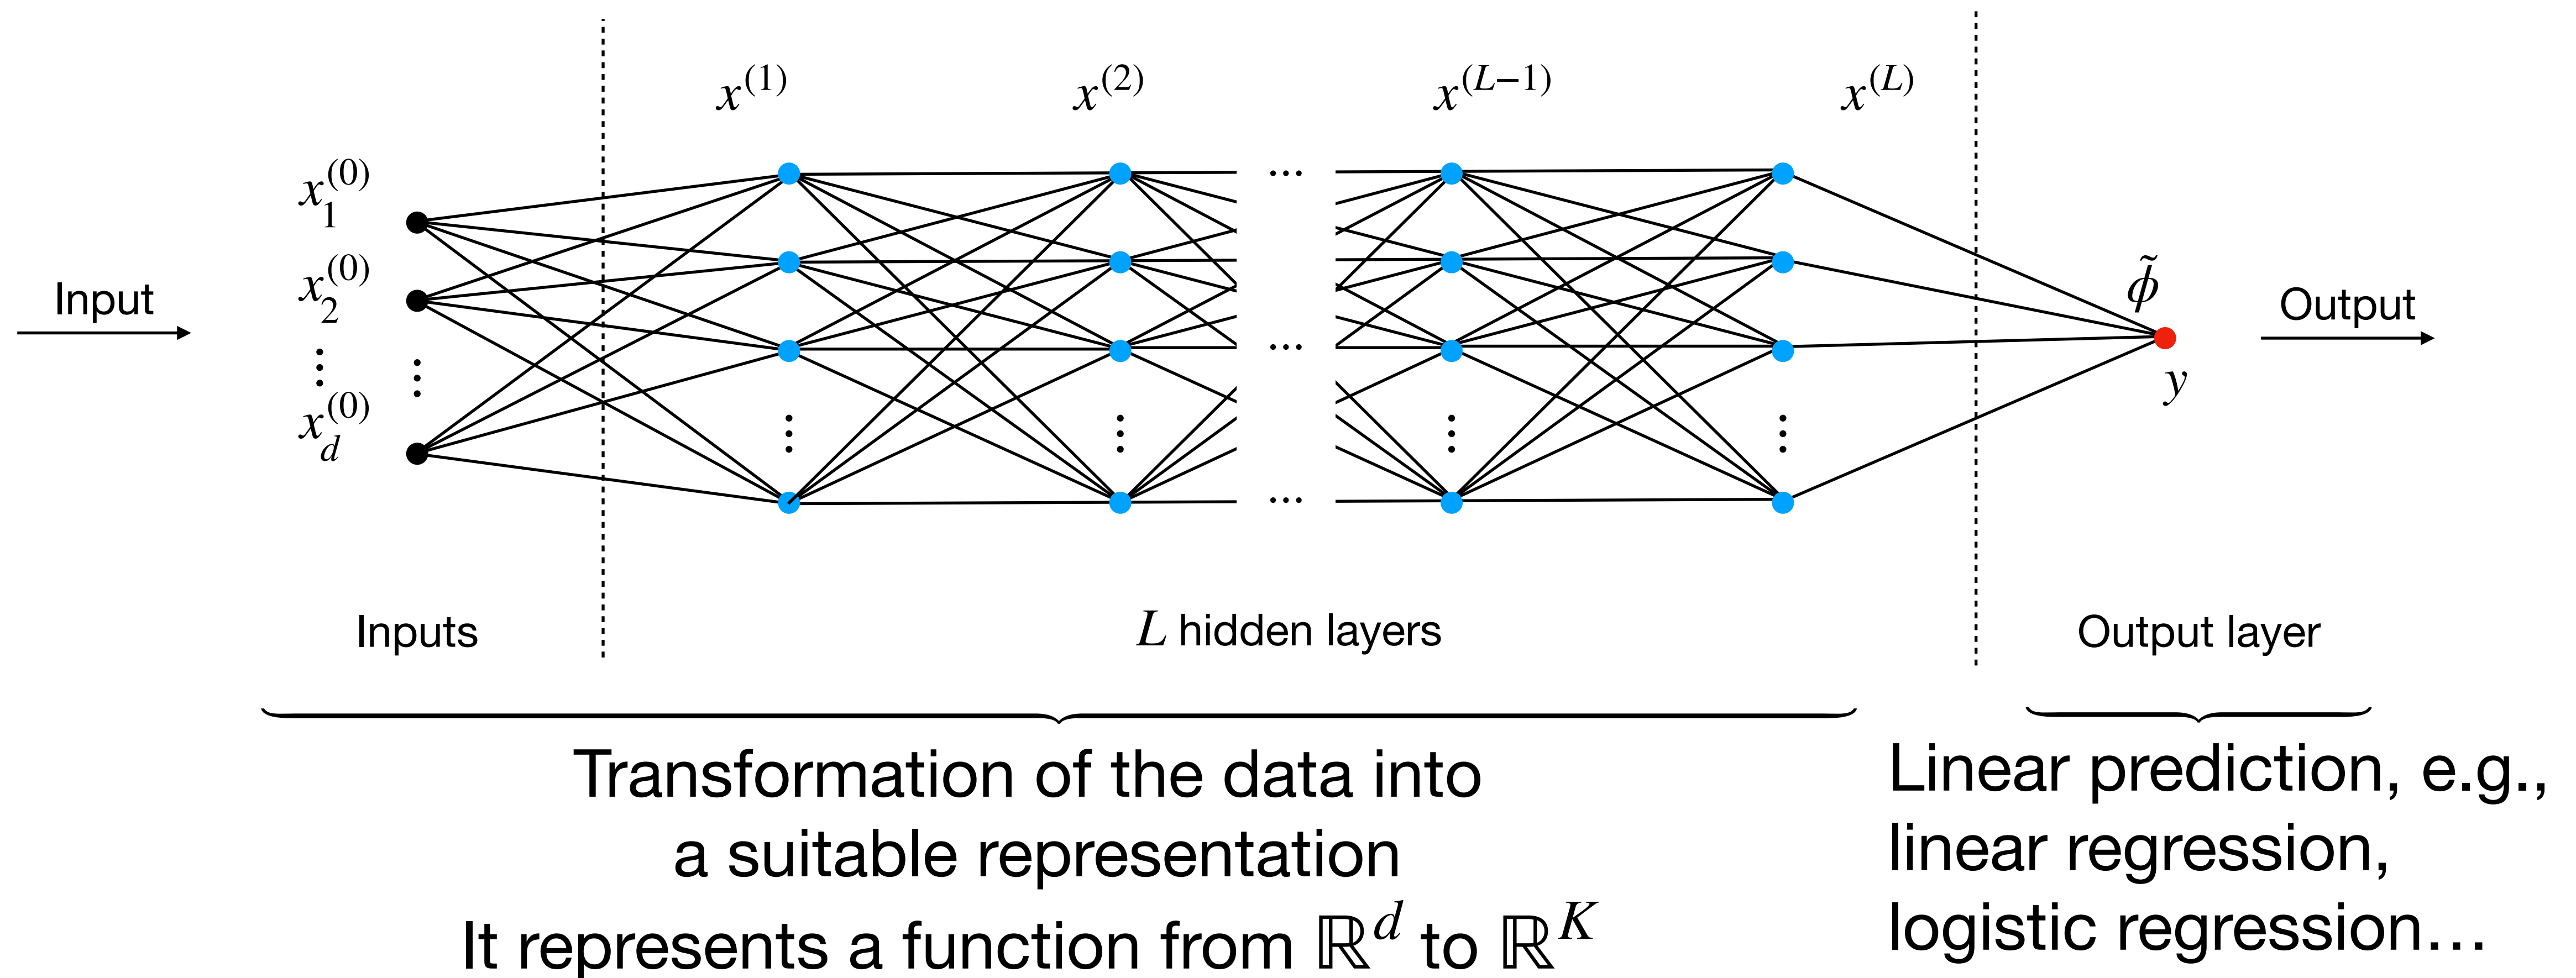
$$x_j^{(l)} = \phi \left(\sum_{i=1}^K x_i^{(l-1)} w_{i,j}^{(l)} + b_j^{(l)} \right)$$

Important: ϕ is non-linear
otherwise we can only
represent linear functions

weight of the edge going from node i
in layer $l - 1$ to node j in layer l

bias term associated with
node j in layer l

The NN transforms the input into a more suitable representation then used to do linear predictions



Representation power

- f smooth (condition on its Fourier coefficients)
- Bounded domain
- Depends on the activation function
- Average approximation in ℓ_2 -norm but also point-wise approximation in ℓ_∞ -norm

Today: How do we train a NN?

Training of NNs

Training loss for a regression problem with $S_{\text{train}} = \{(x_n, y_n)\}_{n=1}^N$:

$$L(f) = \frac{1}{2N} \sum_{n=1}^N (y_n - f(x_n))^2$$

where f is the function represented by a NN with weights $(w_{i,j}^{(l)})$ and biases $(b_i^{(l)})$

Task:

$$\min_{w_{i,j}^{(l)}, b_i^{(l)}} L(f)$$

Remarks:

- Regularization: can be added to avoid overfitting but easy to deal with
- Non convex optimization problem
 - ➡ not guaranteed to converge to a global minimum

Training of NNs with SGD

SGD algorithm:

Sample uniformly n , compute the gradient of $L_n = \frac{1}{2}(y_n - f(x_n))^2$ to update:

$$(w_{i,j}^{(l)})_{t+1} = (w_{i,j}^{(l)})_t - \gamma \frac{\partial}{\partial w_{i,j}^{(l)}} L_n$$

$$(b_i^{(l)})_{t+1} = (b_i^{(l)})_t - \gamma \frac{\partial}{\partial b_i^{(l)}} L_n$$

In Practice: Step size schedule, mini-batch, momentum, Adam

Problem: $O(K^2L)$ parameters: doing chain-rules independently won't be efficient

Solution: Backpropagation algorithm

Compact description of output

The functions implemented by each layer can be written as:

- $x^{(1)} = f^{(1)}(x^{(0)}) := \phi\left((\mathbf{W}^{(1)})^\top x^{(0)} + b^{(1)}\right)$
- $x^{(2)} = f^{(2)}(x^{(1)}) := \phi\left((\mathbf{W}^{(2)})^\top x^{(1)} + b^{(2)}\right)$
- \vdots
 $x^{(l)} = f^{(l)}(x^{(l-1)}) := \phi\left((\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)}\right)$
- \vdots
 $y = f^{(L+1)}(x^{(L)}) := \tilde{\phi}\left((\mathbf{W}^{(L+1)})^\top x^{(L)} + b^{(L+1)}\right)$

The overall function $y = f(x^{(0)})$ is just the composition of these functions:

$$f = f^{(L+1)} \circ f^{(L)} \circ \dots \circ f^{(l)} \circ \dots \circ f^{(2)} \circ f^{(1)}$$

Cost function

Cost function:

$$L = \frac{1}{2N} \sum_{n=1}^N \left(y_n - f^{(L+1)} \circ \dots \circ f^{(2)} \circ f^{(1)}(x_n) \right)^2$$

Remarks:

- The specific form of the loss does not matter
- Function of all weight matrices and bias vectors

Individual loss for SGD:


$$L_n = \frac{1}{2} \left(y_n - f^{(L+1)} \circ \dots \circ f^{(2)} \circ f^{(1)}(x_n) \right)^2$$

Goal: Compute for all (i, j, l)

$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} \quad \text{and} \quad \frac{\partial L_n}{\partial b_i^{(l)}}$$

Naive approach


$$L_n = \frac{1}{2} \left(y_n - \underbrace{f^{(L+1)} \circ \dots \circ f^{(l+1)} \circ \phi \left((\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)} \right)}_{z^{(l)}} \right)^2$$



$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} \quad ?$$

Naive approach

$$L_n = \frac{1}{2} \left(y_n - f^{(L+1)} \circ \dots \circ f^{(l+1)} \circ \phi(z^{(l)}) \right)^2$$



Chain rule:

$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} = \sum_{k=1}^K \frac{\partial L_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}}$$

$$= \frac{\partial L_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial w_{i,j}^{(l)}}$$

$$= \frac{\partial L_n}{\partial z_j^{(l)}} \cdot x_i^{(l-1)}$$

since $\frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}} = 0$ for $k \neq j$

since $z_j^{(l)} = \sum_{k=1}^K w_{k,j}^{(l)} x_k^{(l-1)} + b_j^{(l)}$

We need to compute $\frac{\partial L_n}{\partial z_j^{(l)}}$, $z^{(l)}$ and $x_i^{(l-1)}$


Forward Pass

We can compute $z^{(l)}$ and $x^{(l)}$ by a forward pass in the network:

$$x^{(0)} = x_n \in \mathbb{R}^d$$

$$z^{(l)} = (\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)}$$

$$x^{(l)} = \phi(z^{(l)})$$

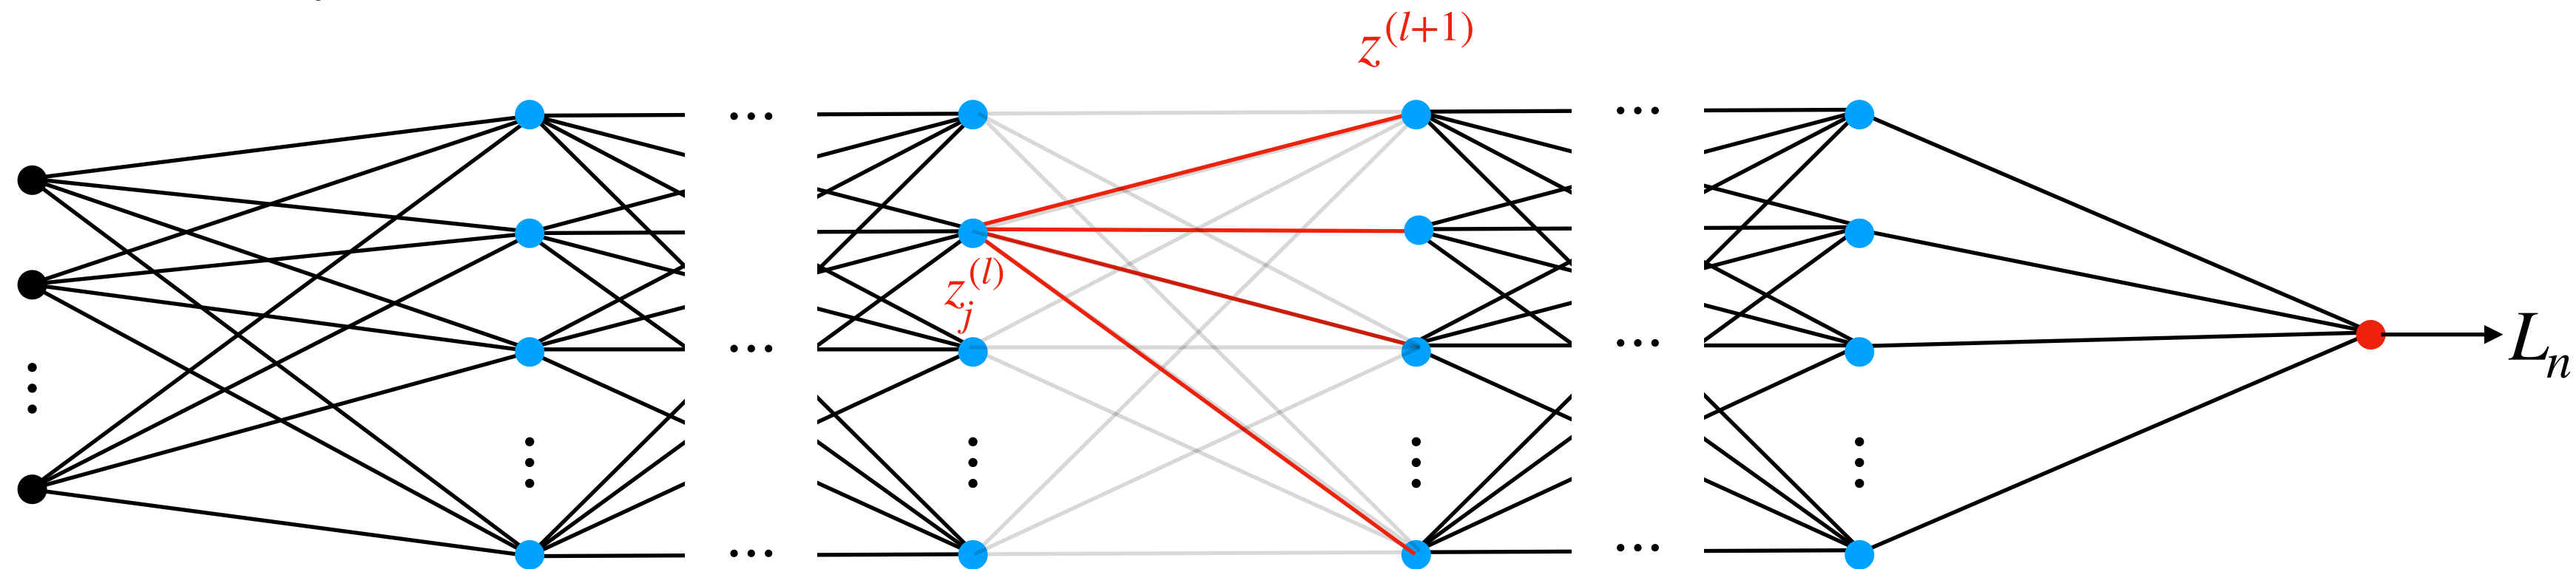


Computational complexity:

➡ one pass over the network $O(K^2L)$

Backward pass (I)

Define $\delta_j^{(l)} = \frac{\partial L_n}{\partial z_j^{(l)}}$



Chain rule:

$$\delta_j^{(l)} = \frac{\partial L_n}{\partial z_j^{(l)}} = \sum_k \frac{\partial L_n}{\partial z_k^{(l+1)}} \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} = \sum_k \delta_k^{(l+1)} \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}}$$

Backward pass (II)

Using $z_k^{(l+1)} = \sum_{i=1}^K w_{i,k}^{(l+1)} x_i^{(l)} + b_k^{(l+1)} = \sum_{i=1}^K w_{i,k}^{(l+1)} \phi(z_i^{(l)}) + b_k^{(l+1)}$

We obtain $\frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} = \phi'(z_j^{(l)}) w_{j,k}^{(l+1)}$

Thus $\delta_j^{(l)} = \sum_k \delta_k^{(l+1)} \phi'(z_j^{(l)}) w_{j,k}^{(l+1)}$

It can be written in vector form:


$$\delta^{(l)} = (\mathbf{W}^{(l+1)} \delta^{(l+1)}) \odot \phi'(z^{(l)})$$

\odot : Hadamard product, i.e.,
pointwise multiplication of vector

Backward pass (III)

Initialization:

$$\begin{aligned}\delta^{(L+1)} &= \frac{\partial}{\partial z^{(L+1)}} \frac{1}{2} (y_n - \tilde{\phi}(z^{(L+1)}))^2 \\ &= (\tilde{\phi}(z^{(L+1)}) - y_n) \tilde{\phi}'(z^{(L+1)})\end{aligned}$$




Compute all the $\delta^{(l)}$ by a backward pass in the network:

$$\begin{aligned}\delta^{(L+1)} &= (\tilde{\phi}(z^{(L+1)}) - y_n) \tilde{\phi}'(z^{(L+1)}) \\ \delta^{(l)} &= (\mathbf{W}^{(l+1)} \delta^{(l+1)}) \odot \phi'(z^{(l)})\end{aligned}$$

Computational complexity: one pass over the network $O(K^2L)$

Derivatives computation



Using that $z_m^{(l)} = \sum_{k=1}^K w_{k,m}^{(l)} x_k^{(l-1)} + b_m^{(l)}$:

- $\frac{\partial L_n}{\partial b_j^{(l)}} = \sum_{k=1}^K \frac{\partial L_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial b_j^{(l)}} = \frac{\partial L_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial b_j^{(l)}} = \delta_j^{(l)}$
- $\frac{\partial L_n}{\partial w_{i,j}^{(l)}} = \sum_{k=1}^K \frac{\partial L_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}} = \frac{\partial L_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} \cdot x_i^{(l-1)}$

Backpropagation algorithm

Forward pass:

$$x^{(0)} = x_n \in \mathbb{R}^d$$

$$z^{(l)} = (\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)}$$

$$x^{(l)} = \phi(z^{(l)})$$

Backward pass:


$$\delta^{(L+1)} = (\tilde{\phi}(z^{(L+1)}) - y_n) \tilde{\phi}'(z^{(L+1)})$$

$$\delta^{(l)} = (\mathbf{W}^{(l+1)} \delta^{(l+1)}) \odot \phi'(z^{(l)})$$

Compute the derivatives:

$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} x_i^{(l-1)}$$

$$\frac{\partial L_n}{\partial b_j^{(l)}} = \delta_j^{(l)}$$



Overall Complexity: $O(K^2L)$

Neural Networks

Popular Activation Functions

Machine Learning Course - CS-433

Nov 9, 2022


Nicolas Flammarion

EPFL

The sigmoid

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$


- Pro: Smooth everywhere
- Cons: $|\sigma'(x)| \lll 1$ for $|x| \ggg 1$ - problem of vanishing gradient



Hyperbolic Tangent


$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2\sigma(2x) - 1$$

- Related to the sigmoid but balanced
- Vanishing gradient problem



Rectified linear unit - RELU


$$(x)_+ = \max\{0, x\}$$



- Pro: no vanishing gradient for $x \geq 0$
- Cons: not differentiable at 0 and the derivative is 0 for negative values

Leaky RELU - LRELU

$$f(x) = \max\{\alpha x, x\}$$




- Correction of the 0 gradient of the RELU

Swish


$$\begin{aligned} f(x) &= x \cdot \sigma(x) \\ &= \frac{x}{1 + e^{-x}} \end{aligned}$$

- Correction of the 0 gradient of the RELU
- Smooth everywhere



Maxout

$$f(x) = \max\{x^\top w_1 + b_1, \dots, x^\top w_k + b_k\}$$



- Generalization of the RELU and the LRELU