

# Neural Networks Training, SGD and Backpropagation

Machine Learning Course - CS-433

Nov 9, 2022

Nicolas Flammarion

**EPFL**

# Recap

# NNs: Key Facts

Supervised learning : we observe some data  $S_{\text{train}} = \{x_i, y_i\}_{i=1}^n \in \mathcal{X} \times \mathcal{Y}$

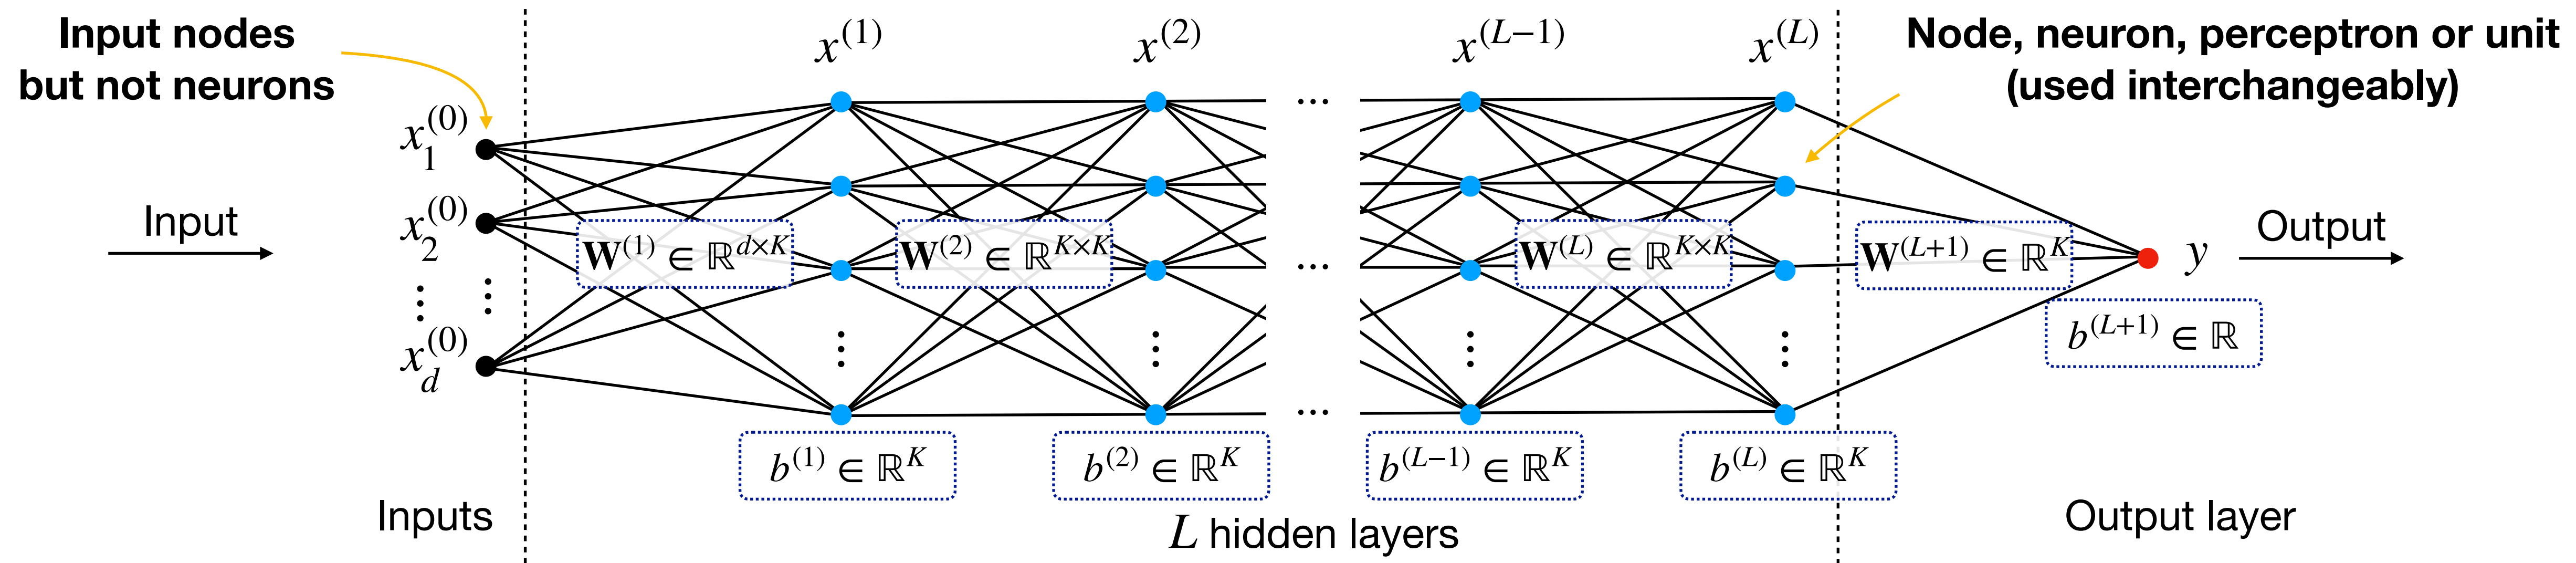
➡ given a new  $x$ , we want to predict its label  $y$

Linear prediction (with augmented features):  $y = f_{\text{Lin}}(x) = \phi(x)^{\top} w$   
Features are given

Prediction with a NN:

$$y = f_{\text{NN}}(x) = f(x)^{\top} w$$
  
Function implemented by the NN parameters: weights and biases  
First layers transform the input into a good representation  
Last layer is performing a linear prediction

# Fully Connected Neural Networks



Assume  $L$  hidden layers with  $K$  neurons each + output layer with single node


Outputs of hidden layer  $l$  given by vector:  $x^{(l)} = f^{(l)}(x^{(l-1)}) := \phi((\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)})$

Elementwise

**Learnable Parameters:** Weight matrices  $\mathbf{W}^{(l)}$  and bias vectors  $b^{(l)}$  for  $1 \leq l \leq L + 1$

- Each column of  $\mathbf{W}^{(l)}$  corresponds to the weights of one perceptron

# Single Neuron View



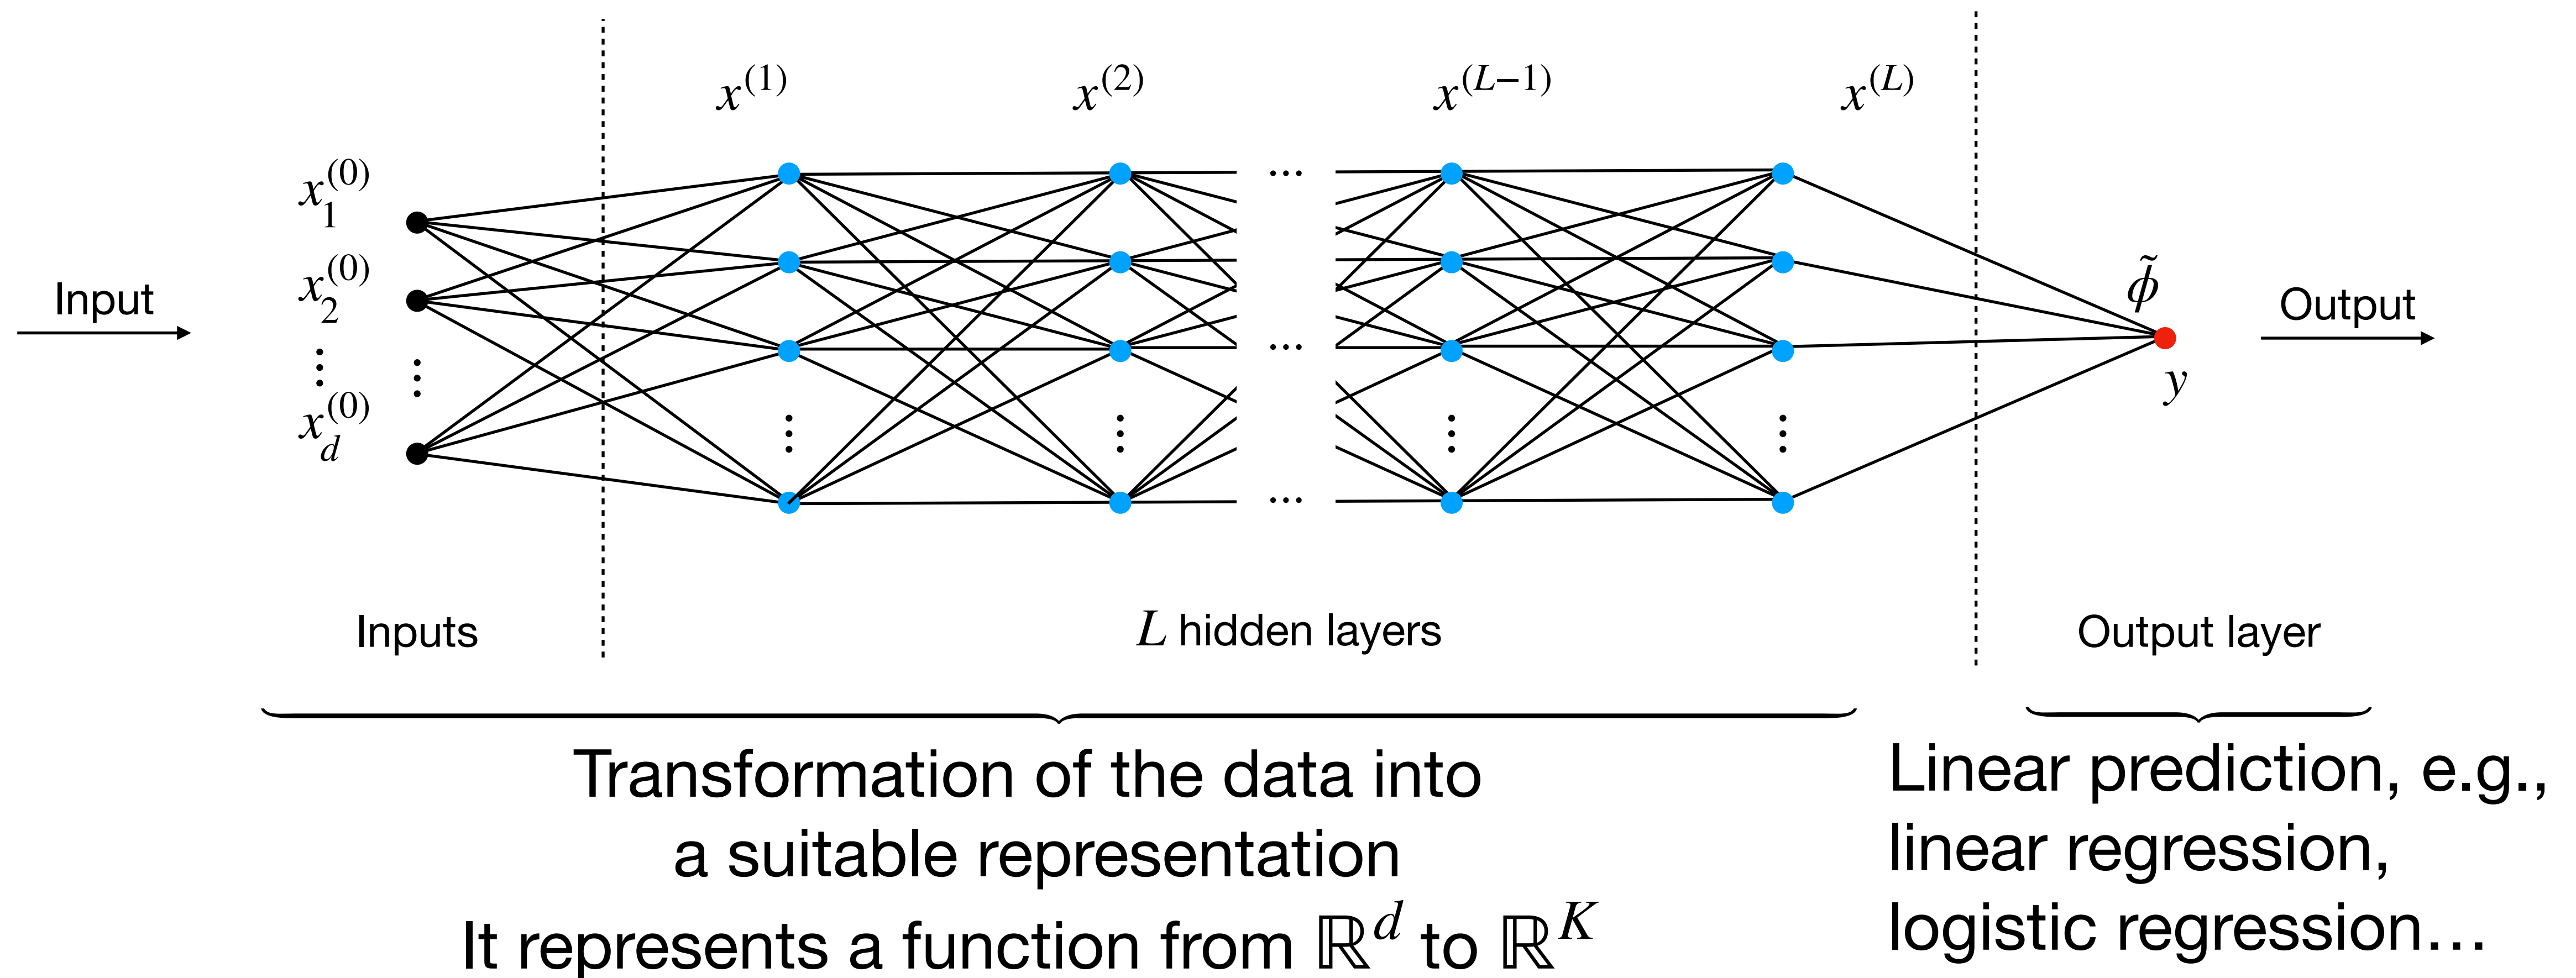
$$x_j^{(l)} = \phi \left( \sum_{i=1}^K x_i^{(l-1)} w_{i,j}^{(l)} + b_j^{(l)} \right)$$

Important:  $\phi$  is non-linear  
otherwise we can only  
represent linear functions

weight of the edge going from node  $i$   
in layer  $l - 1$  to node  $j$  in layer  $l$

bias term associated with  
node  $j$  in layer  $l$

# The NN transforms the input into a more suitable representation then used to do linear predictions



# Representation power

- $f$  smooth (condition on its Fourier coefficients)
- Bounded domain
- Depends on the activation function
- Average approximation in  $\ell_2$ -norm but also point-wise approximation in  $\ell_\infty$ -norm

**Today: How do we train a NN?**



# Training of NNs

Training loss for a regression problem with  $S_{\text{train}} = \{(x_n, y_n)\}_{n=1}^N$ :

$$L(f) = \frac{1}{2N} \sum_{n=1}^N (y_n - f(x_n))^2$$

where  $f$  is the function represented by a NN with weights  $(w_{i,j}^{(l)})$  and biases  $(b_i^{(l)})$

Task:

$$\min_{w_{i,j}^{(l)}, b_i^{(l)}} L(f)$$

Remarks:

- Regularization: can be added to avoid overfitting but easy to deal with
- Non convex optimization problem
  - ➡ not guaranteed to converge to a global minimum

# Training of NNs with SGD

SGD algorithm:

Sample uniformly  $n$ , compute the gradient of  $L_n = \frac{1}{2}(y_n - f(x_n))^2$  to update:

$$(w_{i,j}^{(l)})_{t+1} = (w_{i,j}^{(l)})_t - \gamma \frac{\partial}{\partial w_{i,j}^{(l)}} L_n$$

$$(b_i^{(l)})_{t+1} = (b_i^{(l)})_t - \gamma \frac{\partial}{\partial b_i^{(l)}} L_n$$

In Practice: Step size schedule, momentum, Adam

# Training of NNs with SGD

SGD algorithm:

Sample uniformly  $n$ , compute the gradient of  $L_n = \frac{1}{2}(y_n - f(x_n))^2$  to update:

$$(w_{i,j}^{(l)})_{t+1} = (w_{i,j}^{(l)})_t - \gamma \frac{\partial}{\partial w_{i,j}^{(l)}} L_n$$

$$(b_i^{(l)})_{t+1} = (b_i^{(l)})_t - \gamma \frac{\partial}{\partial b_i^{(l)}} L_n$$

In Practice: Step size schedule, momentum, Adam

Problem:  $O(K^2L)$  parameters

Solution: Backpropagation algorithm

# Compact description of output

The functions implemented by each layer can be written as:

- $x^{(1)} = f^{(1)}(x^{(0)}) := \phi\left((\mathbf{W}^{(1)})^\top x^{(0)} + b^{(1)}\right)$
- $x^{(2)} = f^{(2)}(x^{(1)}) := \phi\left((\mathbf{W}^{(2)})^\top x^{(1)} + b^{(2)}\right)$
- $\vdots$   
 $x^{(l)} = f^{(l)}(x^{(l-1)}) := \phi\left((\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)}\right)$
- $\vdots$   
 $y = f^{(L+1)}(x^{(L)}) := \tilde{\phi}\left((\mathbf{W}^{(L+1)})^\top x^{(L)} + b^{(L+1)}\right)$

The overall function  $y = f(x^{(0)})$  is just the composition of these functions:

$$f = f^{(L+1)} \circ f^{(L)} \circ \dots \circ f^{(l)} \circ \dots \circ f^{(2)} \circ f^{(1)}$$

# Cost function

Cost function:

$$L = \frac{1}{2N} \sum_{n=1}^N \left( y_n - f^{(L+1)} \circ \dots \circ f^{(2)} \circ f^{(1)}(x_n) \right)^2$$

Remarks:

- The specific form of the loss does not matter
- Function of all weight matrices and bias vectors

Individual loss for SGD:


$$L_n = \frac{1}{2} \left( y_n - f^{(L+1)} \circ \dots \circ f^{(2)} \circ f^{(1)}(x_n) \right)^2$$

Goal: Compute for all  $(i, j, l)$

$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} \quad \text{and} \quad \frac{\partial L_n}{\partial b_i^{(l)}}$$

# Naive approach


$$L_n = \frac{1}{2} \left( y_n - \underbrace{f^{(L+1)} \circ \dots \circ f^{(l+1)} \circ \phi \left( (\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)} \right)}_{z^{(l)}} \right)^2$$



$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} \quad ?$$

# Naive approach

$$L_n = \frac{1}{2} \left( y_n - f^{(L+1)} \circ \dots \circ f^{(l+1)} \circ \phi(z^{(l)}) \right)^2$$



Chain rule:

$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} = \sum_{k=1}^K \frac{\partial L_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}}$$

$$= \frac{\partial L_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial w_{i,j}^{(l)}}$$

$$= \frac{\partial L_n}{\partial z_j^{(l)}} \cdot x_i^{(l-1)}$$

since  $\frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}} = 0$  for  $k \neq j$

since  $z_j^{(l)} = \sum_{k=1}^K w_{k,j}^{(l)} x_k^{(l-1)} + b_j^{(l)}$

We need to compute  $\frac{\partial L_n}{\partial z_j^{(l)}}$ ,  $z^{(l)}$  and  $x_i^{(l-1)}$

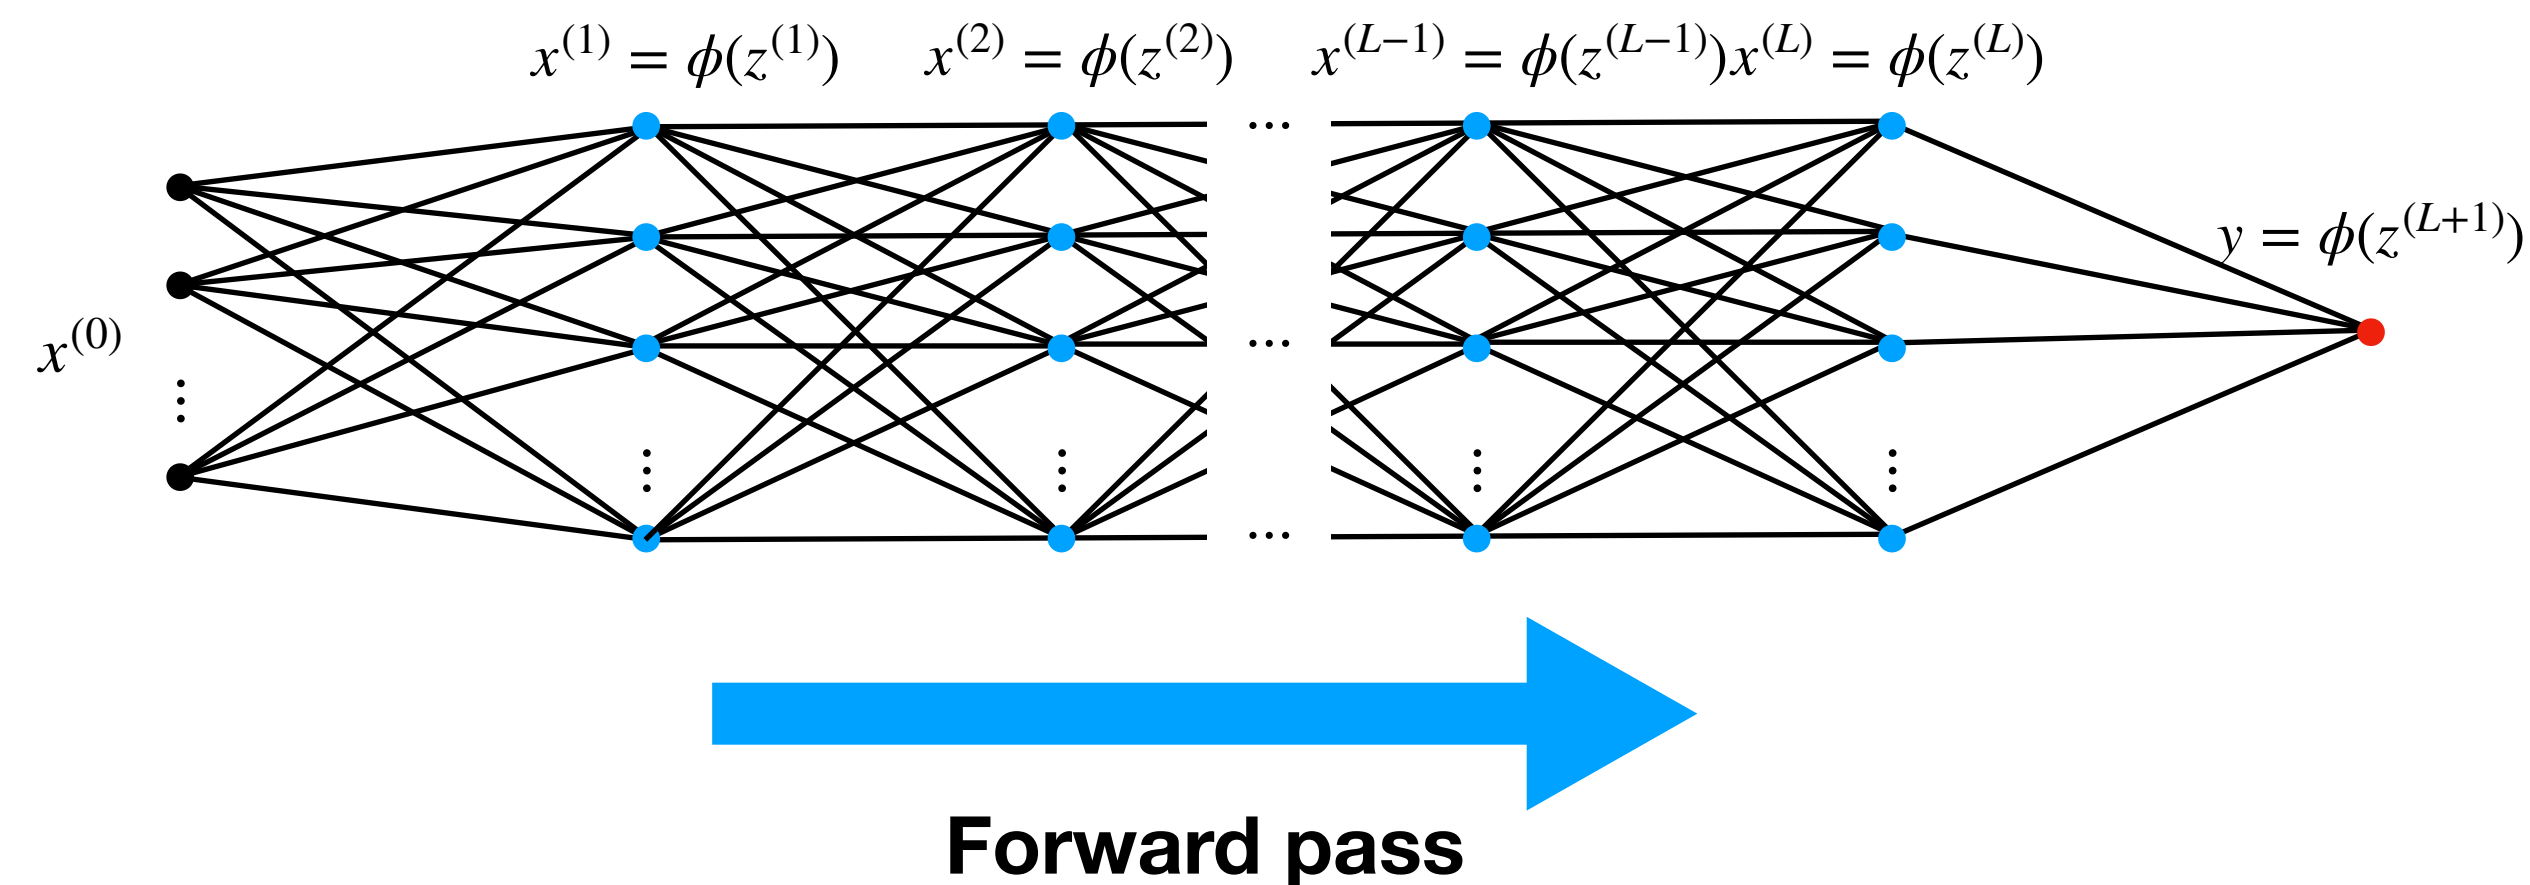
# Forward Pass

We can compute  $z^{(l)}$  and  $x^{(l)}$  by a forward pass in the network:

$$x^{(0)} = x_n \in \mathbb{R}^d$$

$$z^{(l)} = (\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)}$$

$$x^{(l)} = \phi(z^{(l)})$$




Computational complexity:

➡ one pass over the network  $O(K^2L)$



# Backward pass (I)

Define  $\delta_j^{(l)} = \frac{\partial L_n}{\partial z_j^{(l)}}$



Chain rule:

$$\delta_j^{(l)} = \frac{\partial L_n}{\partial z_j^{(l)}} = \sum_k \frac{\partial L_n}{\partial z_k^{(l+1)}} \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} = \sum_k \delta_k^{(l+1)} \frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}}$$

# Backward pass (II)

Using  $z_k^{(l+1)} = \sum_{i=1}^K w_{i,k}^{(l+1)} x_i^{(l)} + b_k^{(l+1)} = \sum_{i=1}^K w_{i,k}^{(l+1)} \phi(z_i^{(l)}) + b_k^{(l+1)}$

We obtain  $\frac{\partial z_k^{(l+1)}}{\partial z_j^{(l)}} = \phi'(z_j^{(l)}) w_{j,k}^{(l+1)}$

Thus  $\delta_j^{(l)} = \sum_k \delta_k^{(l+1)} \phi'(z_j^{(l)}) w_{j,k}^{(l+1)}$

It can be written in vector form:


$$\delta^{(l)} = (\mathbf{W}^{(l+1)} \delta^{(l+1)}) \odot \phi'(z^{(l)})$$

$\odot$ : Hadamard product, i.e.,  
pointwise multiplication of vector

# Backward pass (III)

Initialization:

$$\begin{aligned}\delta^{(L+1)} &= \frac{\partial}{\partial z^{(L+1)}} \frac{1}{2} (y_n - \tilde{\phi}(z^{(L+1)}))^2 \\ &= (\tilde{\phi}(z^{(L+1)}) - y_n) \tilde{\phi}'(z^{(L+1)})\end{aligned}$$



Compute all the  $\delta^{(l)}$  by a backward pass in the network:

$$\begin{aligned}\delta^{(L+1)} &= (\tilde{\phi}(z^{(L+1)}) - y_n) \tilde{\phi}'(z^{(L+1)}) \\ \delta^{(l)} &= (\mathbf{W}^{(l+1)} \delta^{(l+1)}) \odot \phi'(z^{(l)})\end{aligned}$$

Computational complexity: one pass over the network  $O(K^2L)$

# Derivatives computation

Using that  $z_m^{(l)} = \sum_{k=1}^K w_{k,m}^{(l)} x_k^{(l-1)} + b_m^{(l)}$ :

$$\bullet \frac{\partial L_n}{\partial b_j^{(l)}} = \sum_{k=1}^K \frac{\partial L_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial b_j^{(l)}} = \frac{\partial L_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial b_j^{(l)}} = \delta_j^{(l)}$$

$$\bullet \frac{\partial L_n}{\partial w_{i,j}^{(l)}} = \sum_{k=1}^K \frac{\partial L_n}{\partial z_k^{(l)}} \frac{\partial z_k^{(l)}}{\partial w_{i,j}^{(l)}} = \frac{\partial L_n}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} \cdot x_i^{(l-1)}$$

# Backpropagation algorithm

Forward pass:

$$x^{(0)} = x_n \in \mathbb{R}^d$$

$$z^{(l)} = (\mathbf{W}^{(l)})^\top x^{(l-1)} + b^{(l)}$$

$$x^{(l)} = \phi(z^{(l)})$$

Backward pass:

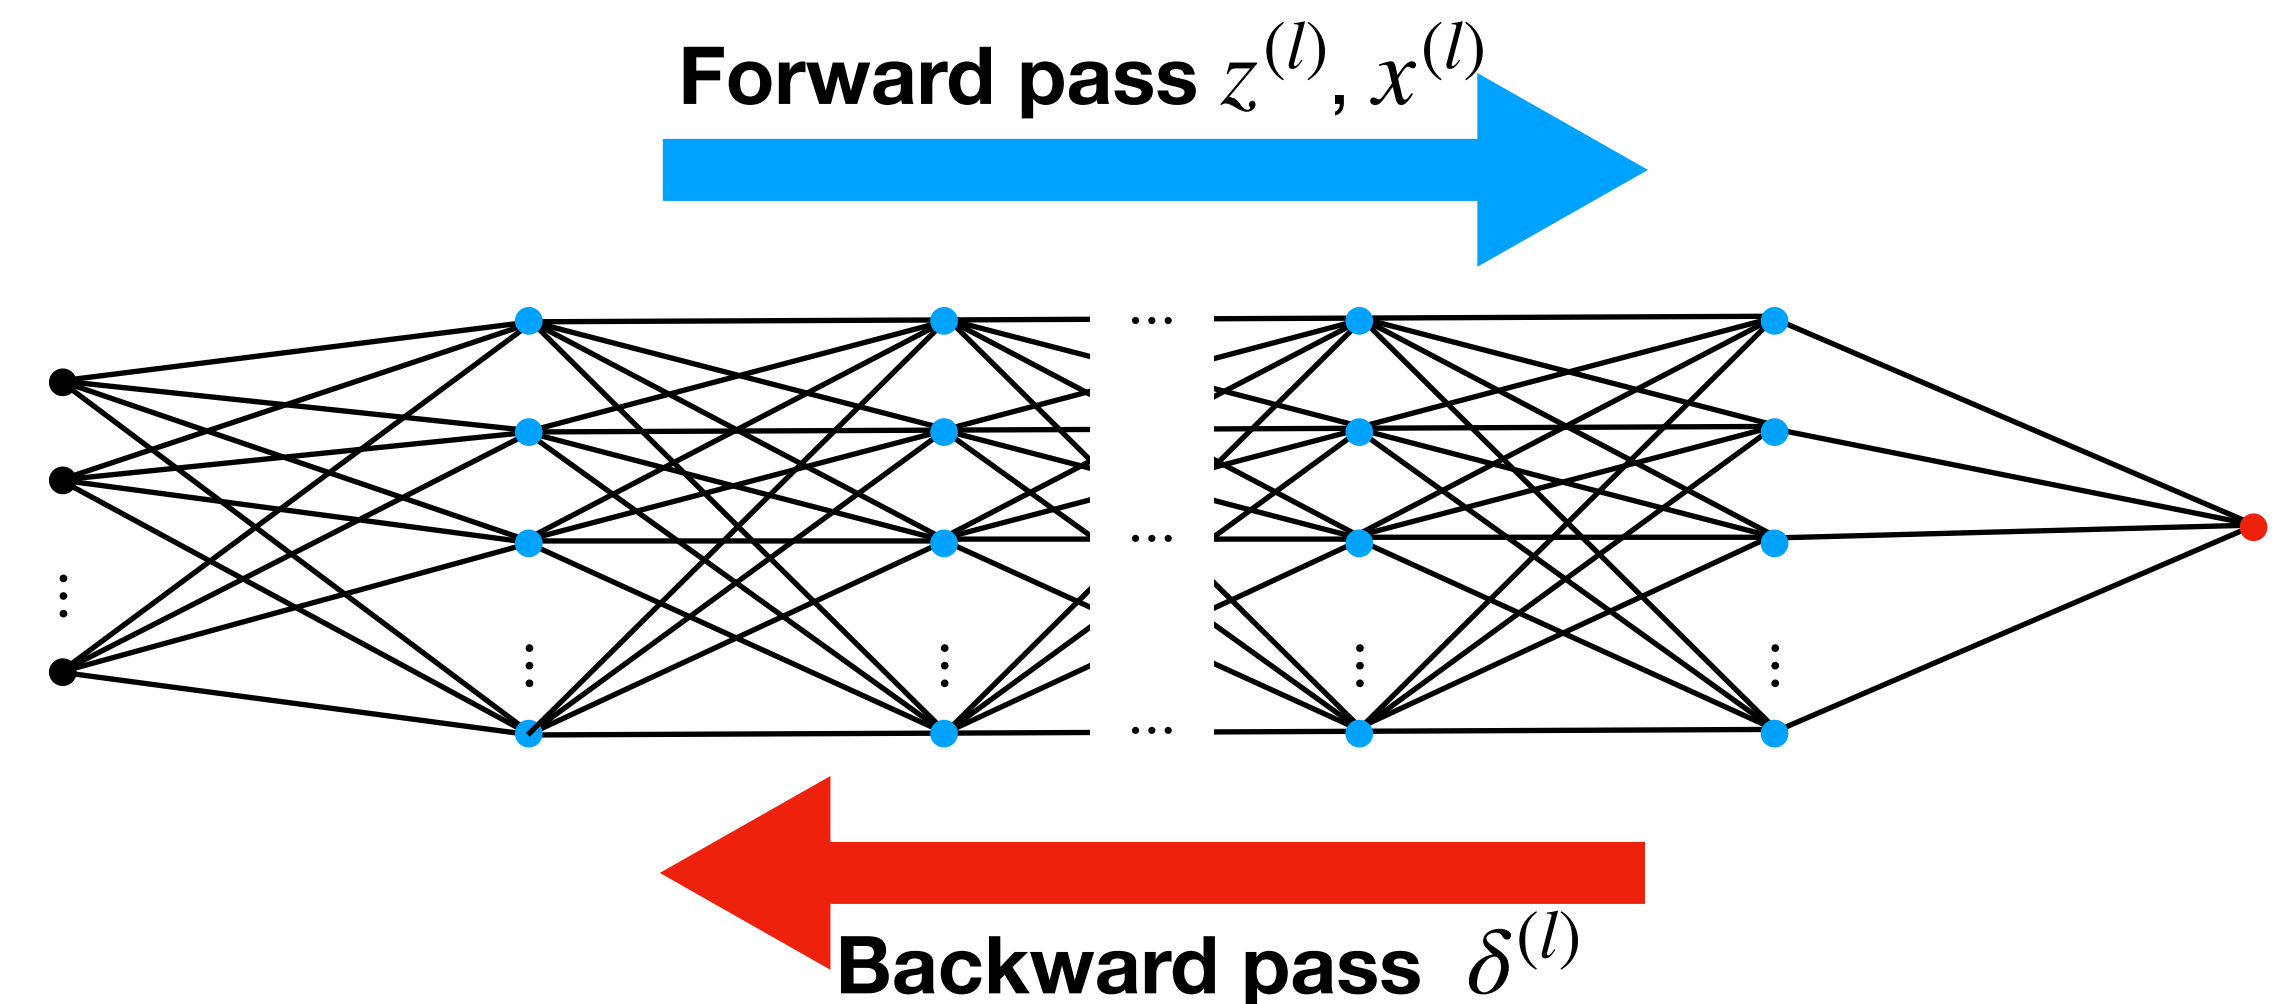
$$\delta^{(L+1)} = (\tilde{\phi}(z^{(L+1)}) - y_n) \tilde{\phi}'(z^{(L+1)})$$

$$\delta^{(l)} = (\mathbf{W}^{(l+1)} \delta^{(l+1)}) \odot \phi'(z^{(l)})$$

Compute the derivatives:

$$\frac{\partial L_n}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} x_i^{(l-1)}$$

$$\frac{\partial L_n}{\partial b_j^{(l)}} = \delta_j^{(l)}$$



Overall Complexity:  $O(K^2L)$

# Neural Networks

# Popular Activation Functions

Machine Learning Course - CS-433


Nov 9, 2022

Nicolas Flammarion

**EPFL**

# The sigmoid

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$




- Pro: Smooth everywhere
- Cons:  $|\sigma'(x)| \lll 1$  for  $|x| \ggg 1$  - problem of vanishing gradient

# Hyperbolic Tangent

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2\sigma(2x) - 1$$


- Related to the sigmoid but balanced
- Vanishing gradient problem





# Rectified linear unit - RELU


$$(x)_+ = \max\{0, x\}$$



- Pro: no vanishing gradient for  $x \geq 0$
- Cons: not differentiable at 0 and the derivative is 0 for negative values

# Leaky RELU - LRELU


$$f(x) = \max\{\alpha x, x\}$$



- Correction of the 0 gradient of the RELU

# Maxout

$$f(x) = \max\{x^\top w_1 + b_1, \dots, x^\top w_k + b_k\}$$



- Generalization of the RELU and the LRELU